


Contents



189 autonomous runs across five targets

Target	Platform	Tests	Passed	App fault	Agent fault
DataGhurhi, survey platform (Deployed)	Web	64	51	4	0
Shobar Khamar, marketplace (Deployed)	Android	45	20	11	14
Google Contacts (Deployed)	Android	50	20	19	10
broken-workshop, (no specification, srs, figma)	Web	10	5	0	4
AcademyBug, (no specification, srs, figma)	Web	20	1	3	16

The two end rows are the whole argument

With a specification

0

agent faults in 64 runs.

Without one

16

agent faults in 20 runs.

The specification is not decoration. It is what lets the agent know where it is going.

USD 0.008

per executed test, measured against the provider's billing record at 10.1 device steps per test.

Comparable published work reports about USD 18 per application, and states that the cost prevented the authors from repeating their own runs.

Limitations and next steps

What we cannot claim yet

- No ground truth: both apps are closed-source, so precision and recall cannot be computed.
- No ablation yet: we have not measured what each component contributes.

What comes next

- An open-source benchmark with known, reproducible bugs (Themis) for bug-detection rate.
- Ablations: findings off, memory off, and a random baseline.

Thank you

Questions are welcome.

With thanks to our supervisor, Dr. A. B. M. Alim Al Islam, and to Samsung R&D Institute Bangladesh.